

# 01\_Dataset\_Analysis

December 14, 2025

## 1 NER Dataset Analysis

### 1.1 1. Import dependencies

```
[1]: # Standard library imports
import json
import sys
from pathlib import Path
from collections import Counter, defaultdict
from typing import Dict, List, Tuple

# Data manipulation and analysis
import pandas as pd
import numpy as np

# Visualization
import matplotlib.pyplot as plt
import seaborn as sns

# Statistical analysis
from scipy import stats

# Add project root to path
project_root = Path.cwd().parent
if str(project_root) not in sys.path:
    sys.path.insert(0, str(project_root))

# Import custom data loader
from src.data import load_processed_data

# Configuration
import warnings
warnings.filterwarnings('ignore')

# Display settings
pd.set_option('display.max_columns', None)
pd.set_option('display.precision', 2)
pd.set_option('display.float_format', '{:.2f}'.format)
```

```

# Plotting configuration
plt.style.use('seaborn-v0_8-whitegrid')
sns.set_context('notebook', font_scale=1.1)
sns.set_palette('deep')

# Figure defaults
plt.rcParams['figure.figsize'] = (14, 6)
plt.rcParams['figure.dpi'] = 100
plt.rcParams['savefig.dpi'] = 300
plt.rcParams['savefig.bbox'] = 'tight'
plt.rcParams['font.size'] = 11
plt.rcParams['axes.labelsize'] = 12
plt.rcParams['axes.titlesize'] = 14
plt.rcParams['xtick.labelsize'] = 10
plt.rcParams['ytick.labelsize'] = 10
plt.rcParams['legend.fontsize'] = 10

```

Python version: 3.12.12

NumPy version: 1.26.4

Pandas version: 2.3.3

Matplotlib version: 3.10.8

Seaborn version: 0.13.2

## 1.2 2. Data Loading and Preprocessing

[2]:

```

# Define data paths
PROJECT_ROOT = Path.cwd().parent
DATA_DIR = PROJECT_ROOT / "data" / "vlps_2018_ner" / "processed"
OUTPUT_DIR = PROJECT_ROOT / "docs"

# Ensure output directory exists
OUTPUT_DIR.mkdir(parents=True, exist_ok=True)

```

[3]:

```

# Load all splits using src.data loader
data_splits = load_processed_data()
train_data = data_splits['train']
dev_data = data_splits['dev']
test_data = data_splits['test']
all_data = train_data + dev_data + test_data

# Basic dataset information
print("Dataset Split Distribution")
print("=" * 50)
print(f"Training set: {len(train_data):5d} records ({len(train_data)/
    len(all_data)*100:5.2f}%)")

```

```

print(f"Development set:{len(dev_data):5d} records ({len(dev_data)/
    len(all_data)*100:5.2f}%)")
print(f"Test set:      {len(test_data):5d} records ({len(test_data)/
    len(all_data)*100:5.2f}%)")
print(f"{'='*50}")
print(f"Total:         {len(all_data):5d} records")

```

Dataset Split Distribution  
=====

Training set: 781 records (60.92%)  
Development set: 260 records (20.28%)  
Test set: 241 records (18.80%)  
=====

Total: 1282 records

[4]: # Display sample record structure

```

sample_record = train_data[0]

print("Sample Record Structure")
print("=" * 80)
print(f"ID: {sample_record['id']}") 
print(f"Topic: {sample_record['topic']}") 
print(f"\nTitle:")
print(f"{sample_record['title'][:100]}...")
print(f"\nText (first 200 characters):")
print(f"{sample_record['text'][:200]}...")
print(f"\nGround Truth Entities:")
for entity_type, entities in sample_record['ground_truth'].items():
    print(f" {entity_type}: {len(entities)} entities")
    if entities:
        print(f"    Example: {entities[0]}")

```

Sample Record Structure  
=====

ID: 23351225  
Topic: Doi song

Title:  
Chuyên đau khi tiêm...

Text (first 200 characters):  
Nỗi đau của người bệnh khi bị tiêm, kéo dài vài giây, vài phút, cùng lấm có thể đến vài ngày. Nhưng nỗi đau của người làm y tế khi đọc, khi nghe những nhận định đó lối, gắt mác phong bì thì còn kéo dài...

Ground Truth Entities:  
person: 0 entities

```
organizations: 1 entities
    Example: Tổ chức y tế Thế giới
address: 0 entities
```

### 1.3 3. Descriptive Statistics

```
[5]: def compute_split_statistics(data: List[Dict], split_name: str) -> Dict:
    """
    Compute comprehensive statistics for a dataset split.

    Parameters:
    -----
    data : List[Dict]
        Dataset records
    split_name : str
        Name of the split

    Returns:
    -----
    Dict
        Dictionary containing computed statistics
    """
    text_lengths = []
    title_lengths = []
    word_counts = []
    entities_per_doc = []

    person_count = 0
    org_count = 0
    addr_count = 0

    for record in data:
        # Text metrics
        text_lengths.append(len(record['text']))
        title_lengths.append(len(record['title']))
        word_counts.append(len(record['text'].split()))

        # Entity counts
        gt = record['ground_truth']
        n_persons = len(gt['person'])
        n_orgs = len(gt['organizations'])
        n_addrs = len(gt['address'])

        person_count += n_persons
        org_count += n_orgs
        addr_count += n_addrs
        entities_per_doc.append(n_persons + n_orgs + n_addrs)
```

```

    return {
        'split': split_name,
        'n_documents': len(data),
        'n_entities_total': person_count + org_count + addr_count,
        'n_persons': person_count,
        'n_organizations': org_count,
        'n_addresses': addr_count,
        'entities_per_doc_mean': np.mean(entities_per_doc),
        'entities_per_doc_std': np.std(entities_per_doc),
        'entities_per_doc_median': np.median(entities_per_doc),
        'text_length_mean': np.mean(text_lengths),
        'text_length_std': np.std(text_lengths),
        'text_length_median': np.median(text_lengths),
        'text_length_p25': np.percentile(text_lengths, 25),
        'text_length_p75': np.percentile(text_lengths, 75),
        'text_length_p95': np.percentile(text_lengths, 95),
        'word_count_mean': np.mean(word_counts),
        'word_count_std': np.std(word_counts),
        'title_length_mean': np.mean(title_lengths),
    }

# Compute statistics for all splits
stats_train = compute_split_statistics(train_data, 'Train')
stats_dev = compute_split_statistics(dev_data, 'Dev')
stats_test = compute_split_statistics(test_data, 'Test')
stats_all = compute_split_statistics(all_data, 'Overall')

# Create summary DataFrame
stats_df = pd.DataFrame([stats_train, stats_dev, stats_test, stats_all])
stats_df = stats_df.set_index('split')

print("\nDataset Statistics Summary")
print("=" * 100)
display(stats_df[['n_documents', 'n_entities_total', 'n_persons',
                  'n_organizations', 'n_addresses', 'entities_per_doc_mean']])

```

## Dataset Statistics Summary

---



---

split	n_documents	n_entities_total	n_persons	n_organizations	\
Train	781	8992	3308	1672	
Dev	260	3032	1067	584	
Test	241	3869	1829	737	
Overall	1282	15893	6204	2993	

	n_addresses	entities_per_doc_mean
split		
Train	4012	11.51
Dev	1381	11.66
Test	1303	16.05
Overall	6696	12.40

### 1.3.1 3.1 Text Length Distribution

```
[6]: # Extract text metrics
text_lengths = np.array([len(r['text']) for r in all_data])
word_counts = np.array([len(r['text']).split() for r in all_data])

# Compute descriptive statistics
text_stats = {
    'Metric': ['Mean', 'Std Dev', 'Median', 'IQR', 'Min', 'Max', '95th Percentile'],
    'Characters': [
        text_lengths.mean(),
        text_lengths.std(),
        np.median(text_lengths),
        np.percentile(text_lengths, 75) - np.percentile(text_lengths, 25),
        text_lengths.min(),
        text_lengths.max(),
        np.percentile(text_lengths, 95)
    ],
    'Words': [
        word_counts.mean(),
        word_counts.std(),
        np.median(word_counts),
        np.percentile(word_counts, 75) - np.percentile(word_counts, 25),
        word_counts.min(),
        word_counts.max(),
        np.percentile(word_counts, 95)
    ]
}

text_stats_df = pd.DataFrame(text_stats)
print("\nText Length Statistics")
print("=" * 60)
display(text_stats_df)
```

Text Length Statistics

---

Metric	Characters	Words
--------	------------	-------

0	Mean	2864.41	627.91
1	Std Dev	1937.59	425.87
2	Median	2468.00	541.00
3	IQR	2248.00	487.25
4	Min	36.00	8.00
5	Max	13848.00	3117.00
6	95th Percentile	6615.10	1443.00

```
[7]: # Visualize text length distributions
fig, axes = plt.subplots(2, 2, figsize=(16, 12))

# Character count distribution
axes[0, 0].hist(text_lengths, bins=60, color='steelblue', alpha=0.7,
                 edgecolor='black', linewidth=0.5)
axes[0, 0].axvline(text_lengths.mean(), color='red', linestyle='--',
                    linewidth=2,
                    label=f'Mean: {text_lengths.mean():.0f}')
axes[0, 0].axvline(np.median(text_lengths), color='green', linestyle='--',
                    linewidth=2,
                    label=f'Median: {np.median(text_lengths):.0f}')
axes[0, 0].set_xlabel('Text Length (characters)')
axes[0, 0].set_ylabel('Frequency')
axes[0, 0].set_title('Distribution of Text Lengths (Characters)')
axes[0, 0].legend()
axes[0, 0].grid(True, alpha=0.3)

# Word count distribution
axes[0, 1].hist(word_counts, bins=60, color='coral', alpha=0.7,
                 edgecolor='black', linewidth=0.5)
axes[0, 1].axvline(word_counts.mean(), color='red', linestyle='--', linewidth=2,
                    label=f'Mean: {word_counts.mean():.0f}')
axes[0, 1].axvline(np.median(word_counts), color='green', linestyle='--',
                    linewidth=2,
                    label=f'Median: {np.median(word_counts):.0f}')
axes[0, 1].set_xlabel('Text Length (words)')
axes[0, 1].set_ylabel('Frequency')
axes[0, 1].set_title('Distribution of Text Lengths (Words)')
axes[0, 1].legend()
axes[0, 1].grid(True, alpha=0.3)

# Box plot for text lengths by split
split_data = [
    [len(r['text'])] for r in train_data],
    [len(r['text'])] for r in dev_data],
    [len(r['text'])] for r in test_data]
]
```


```

bp = axes[1, 0].boxplot(split_data, labels=['Train', 'Dev', 'Test'], u
    ↪patch_artist=True)
for patch in bp['boxes']:
    patch.set_facecolor('lightblue')
axes[1, 0].set_ylabel('Text Length (characters)')
axes[1, 0].set_title('Text Length Distribution by Split')
axes[1, 0].grid(True, alpha=0.3, axis='y')

# Q-Q plot for normality assessment
stats.probplot(text_lengths, dist="norm", plot=axes[1, 1])
axes[1, 1].set_title('Q-Q Plot: Text Length vs Normal Distribution')
axes[1, 1].grid(True, alpha=0.3)

plt.tight_layout()
plt.savefig(OUTPUT_DIR / 'text_length_distribution.png')
plt.show()

```



## 1.4 4. Entity Distribution Analysis

```
[8]: # Collect all entities
all_persons = []
all_orgs = []
all_addrs = []

entities_per_doc = []

for record in all_data:
    gt = record['ground_truth']
    all_persons.extend(gt['person'])
    all_orgs.extend(gt['organizations'])
    all_addrs.extend(gt['address'])

    total = len(gt['person']) + len(gt['organizations']) + len(gt['address'])
    entities_per_doc.append(total)

# Compute entity statistics
entity_stats = pd.DataFrame({
    'Entity Type': ['Person', 'Organization', 'Address', 'Total'],
    'Total Count': [
        len(all_persons),
        len(all_orgs),
        len(all_addrs),
        len(all_persons) + len(all_orgs) + len(all_addrs)
    ],
    'Unique Count': [
        len(set(all_persons)),
        len(set(all_orgs)),
        len(set(all_addrs)),
        len(set(all_persons)) + len(set(all_orgs)) + len(set(all_addrs))
    ]
})
entity_stats['Uniqueness Ratio'] = entity_stats['Unique Count'] / entity_stats['Total Count']
entity_stats['Percentage'] = (entity_stats['Total Count'] / entity_stats.loc[3, 'Total Count']) * 100

print("\nEntity Type Statistics")
print("=" * 80)
display(entity_stats)
```

Entity Type Statistics

	Entity Type	Total Count	Unique Count	Uniqueness Ratio	Percentage
0	Person	6204	4486	0.72	39.04
1	Organization	2993	1815	0.61	18.83
2	Address	6696	2834	0.42	42.13
3	Total	15893	9135	0.57	100.00

```
[9]: # Visualize entity distributions
fig, axes = plt.subplots(2, 2, figsize=(16, 12))

# Entity type distribution across splits (stacked bar)
splits = ['Train', 'Dev', 'Test']
person_counts = [
    sum(len(r['ground_truth']['person']) for r in train_data),
    sum(len(r['ground_truth']['person']) for r in dev_data),
    sum(len(r['ground_truth']['person']) for r in test_data)
]
org_counts = [
    sum(len(r['ground_truth']['organizations']) for r in train_data),
    sum(len(r['ground_truth']['organizations']) for r in dev_data),
    sum(len(r['ground_truth']['organizations']) for r in test_data)
]
addr_counts = [
    sum(len(r['ground_truth']['address']) for r in train_data),
    sum(len(r['ground_truth']['address']) for r in dev_data),
    sum(len(r['ground_truth']['address']) for r in test_data)
]

x = np.arange(len(splits))
width = 0.6

axes[0, 0].bar(x, person_counts, width, label='Person', color='#1f77b4')
axes[0, 0].bar(x, org_counts, width, bottom=person_counts, label='Organization', color='#ff7f0e')
axes[0, 0].bar(x, addr_counts, width,
                bottom=np.array(person_counts) + np.array(org_counts),
                label='Address', color='#2ca02c')
axes[0, 0].set_xlabel('Dataset Split')
axes[0, 0].set_ylabel('Entity Count')
axes[0, 0].set_title('Entity Distribution by Dataset Split')
axes[0, 0].set_xticks(x)
axes[0, 0].set_xticklabels(splits)
axes[0, 0].legend()
axes[0, 0].grid(True, alpha=0.3, axis='y')

# Overall entity type proportions (pie chart)
colors = ['#1f77b4', '#ff7f0e', '#2ca02c']
explode = (0.05, 0.05, 0.05)
```

```

axes[0, 1].pie([len(all_persons), len(all_orgs), len(all_addrs)],
              labels=['Person', 'Organization', 'Address'],
              autopct='%.1f%%',
              colors=colors,
              explode=explode,
              startangle=90)
axes[0, 1].set_title('Overall Entity Type Distribution')


# Entities per document distribution
axes[1, 0].hist(entities_per_doc, bins=40, color='mediumpurple', alpha=0.7,
                 edgecolor='black', linewidth=0.5)
axes[1, 0].axvline(np.mean(entities_per_doc), color='red', linestyle='--',
                   linewidth=2, label=f'Mean: {np.mean(entities_per_doc):.1f}')
axes[1, 0].axvline(np.median(entities_per_doc), color='green', linestyle='--',
                   linewidth=2, label=f'Median: {np.median(entities_per_doc):.
                     <1f}')
axes[1, 0].set_xlabel('Number of Entities per Document')
axes[1, 0].set_ylabel('Frequency')
axes[1, 0].set_title('Distribution of Entity Density')
axes[1, 0].legend()
axes[1, 0].grid(True, alpha=0.3)

# Uniqueness comparison
entity_types = ['Person', 'Organization', 'Address']
total_counts = [len(all_persons), len(all_orgs), len(all_addrs)]
unique_counts = [len(set(all_persons)), len(set(all_orgs)), len(set(all_addrs))]

x_pos = np.arange(len(entity_types))
axes[1, 1].bar(x_pos - 0.2, total_counts, 0.4, label='Total', color='skyblue')
axes[1, 1].bar(x_pos + 0.2, unique_counts, 0.4, label='Unique', color='orange')
axes[1, 1].set_xlabel('Entity Type')
axes[1, 1].set_ylabel('Count')
axes[1, 1].set_title('Total vs Unique Entity Counts')
axes[1, 1].set_xticks(x_pos)
axes[1, 1].set_xticklabels(entity_types)
axes[1, 1].legend()
axes[1, 1].grid(True, alpha=0.3, axis='y')

plt.tight_layout()
plt.savefig(OUTPUT_DIR / 'entity_distribution.png')
plt.show()

```



## 1.5 5. Topic Analysis

```
[10]: # Analyze topic distribution
topics = Counter([r['topic'] for r in all_data])
topic_df = pd.DataFrame([
    {'Topic': topic, 'Count': count, 'Percentage': (count / len(all_data)) * 100}
    for topic, count in topics.most_common()])
print(f"\nTopic Distribution (Total: {len(topics)} unique topics)")
print("=" * 70)
display(topic_df.head(15))

# Statistical summary
print(f"\nTopic Distribution Statistics:")
print(f"Mean documents per topic: {topic_df['Count'].mean():.2f}")
print(f"Std deviation: {topic_df['Count'].std():.2f}")
print(f"Median documents per topic: {topic_df['Count'].median():.2f}")
```

```
print(f"Coefficient of variation: {(topic_df['Count'].std() / topic_df['Count'].mean()):.2f}")
```

Topic Distribution (Total: 27 unique topics)

---

	Topic	Count	Percentage
0	Van hoa	129	10.06
1	Xa hoi	124	9.67
2	Kinh te	120	9.36
3	Giai tri	118	9.20
4	Phap luat	115	8.97
5	The thao	110	8.58
6	Giao duc	109	8.50
7	The gioi	109	8.50
8	KH - CN	100	7.80
9	Doi song	93	7.25
10	Doi Song	25	1.95
11	Xa hoi Thoi su	17	1.33
12	Kinh te Tai chinh	15	1.17
13	The thao Bong da quoc te	14	1.09
14	Kinh te Kinh doanh	13	1.01

Topic Distribution Statistics:

Mean documents per topic: 47.48

Std deviation: 51.63

Median documents per topic: 14.00

Coefficient of variation: 1.09

```
[11]: # Visualize topic distribution
fig, axes = plt.subplots(1, 2, figsize=(18, 7))

# Top topics horizontal bar chart
top_topics = topic_df.head(15).sort_values('Count')
axes[0].barh(range(len(top_topics)), top_topics['Count'], color='steelblue')
axes[0].set_yticks(range(len(top_topics)))
axes[0].set_yticklabels(top_topics['Topic'])
axes[0].set_xlabel('Number of Documents')
axes[0].set_title('Top 15 Topics by Document Count')
axes[0].grid(True, alpha=0.3, axis='x')

# Add count labels
for i, (idx, row) in enumerate(top_topics.iterrows()):
    axes[0].text(row['Count'] + 1, i, str(int(row['Count'])),
                 va='center', fontsize=9)
```


```

# Topic distribution across splits
topic_split_data = []
for topic in topic_df.head(10)['Topic']:
    topic_split_data.append({
        'Topic': topic,
        'Train': sum(1 for r in train_data if r['topic'] == topic),
        'Dev': sum(1 for r in dev_data if r['topic'] == topic),
        'Test': sum(1 for r in test_data if r['topic'] == topic)
    })

topic_split_df = pd.DataFrame(topic_split_data)
topic_split_df.set_index('Topic')[['Train', 'Dev', 'Test']].plot(
    kind='bar', ax=axes[1], width=0.8
)
axes[1].set_xlabel('Topic')
axes[1].set_ylabel('Number of Documents')
axes[1].set_title('Top 10 Topics Distribution Across Splits')
axes[1].legend(title='Split')
axes[1].tick_params(axis='x', rotation=45)
axes[1].grid(True, alpha=0.3, axis='y')

plt.tight_layout()
plt.savefig(OUTPUT_DIR / 'topic_distribution.png')
plt.show()

```



## 1.6 6. Entity Frequency Analysis

```
[12]: # Count entity occurrences
person_counter = Counter(all_persons)
org_counter = Counter(all_orgs)
addr_counter = Counter(all_addrs)
```

```

# Get top entities
TOP_N = 20

top_persons_df = pd.DataFrame(
    person_counter.most_common(TOP_N),
    columns=['Entity', 'Frequency']
)
top_persons_df['Type'] = 'Person'

top_orgs_df = pd.DataFrame(
    org_counter.most_common(TOP_N),
    columns=['Entity', 'Frequency']
)
top_orgs_df['Type'] = 'Organization'

top_addrs_df = pd.DataFrame(
    addr_counter.most_common(TOP_N),
    columns=['Entity', 'Frequency']
)
top_addrs_df['Type'] = 'Address'

print(f"\nTop {TOP_N} Most Frequent Entities by Type")
print("==" * 80)

print("\nPersons:")
display(top_persons_df.head(10))

print("\nOrganizations:")
display(top_orgs_df.head(10))

print("\nAddresses:")
display(top_addrs_df.head(10))

```

Top 20 Most Frequent Entities by Type

---

Persons:

	Entity	Frequency	Type
0	Donald Trump	32	Person
1	Trump	20	Person
2	Hà Văn Thắm	16	Person
3	Vũ Hồng Việt	15	Person
4	Sơn	14	Person
5	Nguyễn Hoàng	14	Person
6	Thảo	13	Person
7	Lâm	12	Person

8	Nguyễn Thị Nghĩa	12	Person
9	Messi	12	Person

Organizations:

	Entity	Frequency	Type
0	TTXVN	34	Organization
1	Reuters	27	Organization
2	Vietnam+	27	Organization
3	AFP	26	Organization
4	Apple	24	Organization
5	Bộ Chính trị	18	Organization
6	Hồi giáo	18	Organization
7	Bộ Công Thương	17	Organization
8	VOV	16	Organization
9	Chelsea	16	Organization

Addresses:

	Entity	Frequency	Type
0	Việt Nam	328	Address
1	Mỹ	177	Address
2	Hà Nội	174	Address
3	Trung Quốc	112	Address
4	TP.HCM	75	Address
5	Anh	68	Address
6	Nga	68	Address
7	Nhật Bản	57	Address
8	Hàn Quốc	56	Address
9	châu Âu	51	Address

```
[13]: # Visualize top entities
fig, axes = plt.subplots(3, 1, figsize=(14, 18))

# Top persons
top_persons = top_persons_df.head(15).sort_values('Frequency')
axes[0].barh(range(len(top_persons)), top_persons['Frequency'], color="#1f77b4")
axes[0].set_yticks(range(len(top_persons)))
axes[0].set_yticklabels(top_persons['Entity'])
axes[0].set_xlabel('Frequency')
axes[0].set_title('Top 15 Most Mentioned Persons')
axes[0].grid(True, alpha=0.3, axis='x')
for i, (idx, row) in enumerate(top_persons.iterrows()):
    axes[0].text(row['Frequency'] + 0.5, i, str(int(row['Frequency'])),
                 va='center', fontsize=9)

# Top organizations
```

```


top_orgs = top_orgs_df.head(15).sort_values('Frequency')
axes[1].barh(range(len(top_orgs)), top_orgs['Frequency'], color='#ff7f0e')
axes[1].set_yticks(range(len(top_orgs)))
axes[1].set_yticklabels(top_orgs['Entity'])
axes[1].set_xlabel('Frequency')
axes[1].set_title('Top 15 Most Mentioned Organizations')
axes[1].grid(True, alpha=0.3, axis='x')
for i, (idx, row) in enumerate(top_orgs.iterrows()):
    axes[1].text(row['Frequency'] + 0.5, i, str(int(row['Frequency'])),
                 va='center', fontsize=9)

# Top addresses
top_addrs = top_addrs_df.head(15).sort_values('Frequency')
axes[2].barh(range(len(top_addrs)), top_addrs['Frequency'], color='#2ca02c')
axes[2].set_yticks(range(len(top_addrs)))
axes[2].set_yticklabels(top_addrs['Entity'])
axes[2].set_xlabel('Frequency')
axes[2].set_title('Top 15 Most Mentioned Addresses')
axes[2].grid(True, alpha=0.3, axis='x')
for i, (idx, row) in enumerate(top_addrs.iterrows()):
    axes[2].text(row['Frequency'] + 1, i, str(int(row['Frequency'])),
                 va='center', fontsize=9)


plt.tight_layout()
plt.savefig(OUTPUT_DIR / 'top_entities.png')
plt.show()

```


Top 15 Most Mentioned Persons



Top 15 Most Mentioned Organizations



Top 15 Most Mentioned Addresses



## 1.7 7. Entity Co-occurrence Analysis

```
[14]: # Analyze entity co-occurrence patterns
cooccurrence = defaultdict(int)

for record in all_data:
    gt = record['ground_truth']
    has_person = len(gt['person']) > 0
    has_org = len(gt['organizations']) > 0
    has_addr = len(gt['address']) > 0

    if has_person and has_org and has_addr:
        cooccurrence['All Three'] += 1
    elif has_person and has_org:
        cooccurrence['Person + Organization'] += 1
    elif has_person and has_addr:
        cooccurrence['Person + Address'] += 1
    elif has_org and has_addr:
        cooccurrence['Organization + Address'] += 1
    elif has_person:
        cooccurrence['Person Only'] += 1
    elif has_org:
        cooccurrence['Organization Only'] += 1
    elif has_addr:
        cooccurrence['Address Only'] += 1
    else:
        cooccurrence['No Entities'] += 1

# Create DataFrame
cooccurrence_df = pd.DataFrame([
    {'Pattern': k, 'Count': v, 'Percentage': (v / len(all_data)) * 100}
    for k, v in cooccurrence.items()])
]).sort_values('Count', ascending=False)

print("\nEntity Co-occurrence Patterns")
print("=" * 70)
display(cooccurrence_df)
```

Entity Co-occurrence Patterns

=====

	Pattern	Count	Percentage
2	All Three	676	52.73
6	Person + Address	271	21.14
5	Organization + Address	77	6.01
1	Person + Organization	62	4.84
3	Person Only	61	4.76

7	Address Only	53	4.13
4	No Entities	52	4.06
0	Organization Only	30	2.34

```
[15]: # Visualize co-occurrence patterns
fig, ax = plt.subplots(figsize=(12, 7))

colors = plt.cm.Set3(range(len(cooccurrence_df)))
bars = ax.bar(range(len(cooccurrence_df)), cooccurrence_df['Count'],
              color=colors, edgecolor='black', linewidth=1.2)
ax.set_xticks(range(len(cooccurrence_df)))
ax.set_xticklabels(cooccurrence_df['Pattern'], rotation=45, ha='right')
ax.set_xlabel('Entity Combination Pattern')
ax.set_ylabel('Number of Documents')
ax.set_title('Entity Co-occurrence Patterns in Documents')
ax.grid(True, alpha=0.3, axis='y')

# Add value labels
for i, (idx, row) in enumerate(cooccurrence_df.iterrows()):
    height = bars[i].get_height()
    ax.text(bars[i].get_x() + bars[i].get_width()/2., height,
            f'{int(row["Count"])}\n{row["Percentage"]:.1f}%',
            ha='center', va='bottom', fontsize=9)

plt.tight_layout()
plt.savefig(OUTPUT_DIR / 'cooccurrence_patterns.png')
plt.show()
```

